

Object Recognition

The objective of this lab is very simple, to recognize objects in images. You will be working with a well-known dataset called CIFAR-10.

You can learn more about this dataset and download it here:

<https://www.cs.toronto.edu/~kriz/cifar.html> (<https://www.cs.toronto.edu/~kriz/cifar.html>)

In the webpage above, they also included a few publications based on CIFAR-10 data, which showed some amazing accuracies. The worst network on the page (a shallow convolutional neural network) can classify images with roughly 75% accuracy.

1. Write a function to load data

The dataset webpage in the previous section also provide a simple way to load data from your harddrive using pickle. You may use their function for this exercise.

Construct two numpy arrays for train images and train labels from data_batch_1 to data_batch_5. Then, construct two numpy arrays for test images, and test labels from test batch file. The original image size is 32 x 32 x 3. You may flatten the arrays so the final arrays are of size 1 x 3072.

In [407]:

```
from matplotlib import pyplot
from keras.datasets import cifar10
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

(X_train, y_train), (X_test, y_test) = cifar10.load_data()
```

In [408]:

```
X_train = X_train.reshape(50000, 3072)
y_train = y_train.reshape(-1)
print(X_train.shape)
print(y_train.shape)

X_test = X_test.reshape(10000, 3072)
y_test = y_test.reshape(-1)
print(X_test.shape)
print(y_test.shape)
```

```
(50000, 3072)
(50000,)
(10000, 3072)
(10000,)
```

In [409]:

```
X_train = X_train / 255
X_test = X_test / 255
```

2. Classify Dogs v.s. Cats

Let's start simple by creating logistic regression model to classify images. We will select only two classes of images for this exercise.

1. From 50,000 train images and 10,000 test images, we want to reduce the data size. Write code to filter only dog images (label = 3) and cat images (label = 5).
2. Create a logistic regression model to classify cats and dogs. Report your accuracy.

In [410]:

```
dog_images_train = []
for i in range (0, 50000):
    if labels[y_train[i]] == 'dog':
        dog_images_train.append(X_train[i])
print(len(dog_images_train))

dog_images_test = []
for i in range (0, 10000):
    if labels[y_test[i]] == 'dog':
        dog_images_test.append(X_test[i])
print(len(dog_images_test))
```

5000

1000

In [411]:

```
dog_labels_train = [3] * 5000
dog_labels_test = [3] * 1000
```

In [412]:

```
cat_images_train = []
for i in range (0, 50000):
    if labels[y_train[i]] == 'cat':
        cat_images_train.append(X_train[i])
print(len(dog_images_train))

cat_images_test = []
for i in range (0, 10000):
    if labels[y_test[i]] == 'cat':
        cat_images_test.append(X_test[i])
print(len(dog_images_test))
```

5000

1000

In [413]:

```
cat_labels_train = [5] * 5000  
cat_labels_test = [5] * 1000
```

In [414]:

```
dog_cat_images_train = np.array(dog_images_train + cat_images_train)  
dog_cat_labels_train = np.array(dog_labels_train + cat_labels_train)  
  
dog_cat_images_test = np.array(dog_images_test + cat_images_test)  
dog_cat_labels_test = np.array(dog_labels_test + cat_labels_test)
```

In [415]:

```
from sklearn.linear_model import LogisticRegression  
LR = LogisticRegression(random_state = 0, max_iter = 2000)  
LR.fit(dog_cat_images_train, dog_cat_labels)
```

Out[415]:

```
LogisticRegression(max_iter=2000, random_state=0)
```

In [416]:

```
y_pred = LR.predict(dog_cat_images_test)
```

In [417]:

```
from sklearn.metrics import accuracy_score  
accuracy_score(dog_cat_labels_test, y_pred)
```

Out[417]:

```
0.5805
```

3. The Real Challenge

The majority of your score for this lab will come from this real challenge. You are going to construct a neural network model to classify 10 classes of images from CIFAR-10 dataset. You will get half the credits for this one if you complete the assignment, and will get another half if you can exceed the target accuracy of 75%. (You may use any combination of sklearn, opencv, or tensorflow to do this exercise).

Design at least 3 variants of neural network models. Each model should have different architectures. (Do not vary just a few parameters, the architecture of the network must change in each model). In your notebook, explain your experiments in details and display the accuracy score for each experiment.

In [512]:

```
# first experiment using CNN
import numpy
from keras.datasets import cifar10
from keras.models import Sequential
from keras.layers import Dense,Dropout,Flatten,Conv2D,MaxPooling2D
from keras.constraints import maxnorm
from keras.optimizers import SGD
from keras.utils import np_utils
from keras import backend
import keras as K
```

In [513]:

```
(X_train, y_train), (X_test, y_test) = cifar10.load_data()
```

In [514]:

```
#normalize the inputs from 0-255 to 0-1 ,as pixels have value in the range 0-255 hence we n
X_train = X_train.astype('float32')
X_test = X_test.astype('float32')
X_train = X_train / 255.0
X_test = X_test / 255.0
```

In [515]:

```
# one hot encode outputs for better prediction
y_train = np_utils.to_categorical(y_train)
y_test = np_utils.to_categorical(y_test)
num_classes = y_test.shape[1]
```

In [516]:

```
# Create the model using sequentia as it allows to create the model layer-by-layer
model = Sequential()
```

In []:

```
# add all the required layers for CNN
model.add(Conv2D(32, (3, 3), input_shape=(32,32,3), activation='relu', padding='same'))
model.add(Dropout(0.2))
model.add(Conv2D(32, (3, 3), activation='relu', padding='same'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Conv2D(64, (3, 3), activation='relu', padding='same'))
model.add(Dropout(0.2))
model.add(Conv2D(64, (3, 3), activation='relu', padding='same'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Conv2D(128, (3, 3), activation='relu', padding='same'))
model.add(Dropout(0.2))
model.add(Conv2D(128, (3, 3), activation='relu', padding='same'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Flatten())
model.add(Dropout(0.2))
model.add(Dense(1024, activation='relu', kernel_constraint=maxnorm(3)))
model.add(Dropout(0.2))
model.add(Dense(512, activation='relu', kernel_constraint=maxnorm(3)))
model.add(Dropout(0.2))
model.add(Dense(num_classes, activation='softmax'))
```

In [542]:

```
model.summary()
```

Model: "sequential_17"

Layer (type)	Output Shape	Param #
=====		
conv2d_34 (Conv2D)	(None, 32, 32, 32)	896
dropout_24 (Dropout)	(None, 32, 32, 32)	0
conv2d_35 (Conv2D)	(None, 32, 32, 32)	9248
max_pooling2d_20 (MaxPooling)	(None, 16, 16, 32)	0
conv2d_36 (Conv2D)	(None, 16, 16, 64)	18496
dropout_25 (Dropout)	(None, 16, 16, 64)	0
conv2d_37 (Conv2D)	(None, 16, 16, 64)	36928
max_pooling2d_21 (MaxPooling)	(None, 8, 8, 64)	0
conv2d_38 (Conv2D)	(None, 8, 8, 128)	73856
dropout_26 (Dropout)	(None, 8, 8, 128)	0
conv2d_39 (Conv2D)	(None, 8, 8, 128)	147584
max_pooling2d_22 (MaxPooling)	(None, 4, 4, 128)	0
flatten_8 (Flatten)	(None, 2048)	0
dropout_27 (Dropout)	(None, 2048)	0
dense_16 (Dense)	(None, 1024)	2098176
dropout_28 (Dropout)	(None, 1024)	0
dense_17 (Dense)	(None, 512)	524800
dropout_29 (Dropout)	(None, 512)	0
dense_18 (Dense)	(None, 10)	5130
=====		
Total params: 2,915,114		
Trainable params: 2,915,114		
Non-trainable params: 0		

In [534]:

```
# Compile model
epochs = 10
lr = 0.01
decay = lr/epochs
sgd = SGD(lr = lr, momentum=0.9, decay=decay, nesterov=False,)
model.compile(loss='categorical_crossentropy', optimizer=sgd, metrics=['accuracy'])
model.fit(X_train, y_train, validation_data=(X_test, y_test), epochs=40, batch_size=32)
```

Epoch 25/40

1563/1563 [=====] - 151s 97ms/step - loss: 0.3487
- accuracy: 0.8763 - val_loss: 0.6233 - val_accuracy: 0.7918

Epoch 26/40

1563/1563 [=====] - 150s 96ms/step - loss: 0.3385
- accuracy: 0.8792 - val_loss: 0.6207 - val_accuracy: 0.7947

Epoch 27/40

1563/1563 [=====] - 150s 96ms/step - loss: 0.3388
- accuracy: 0.8795 - val_loss: 0.6176 - val_accuracy: 0.7962

Epoch 28/40

1563/1563 [=====] - 151s 97ms/step - loss: 0.3369
- accuracy: 0.8795 - val_loss: 0.6216 - val_accuracy: 0.7952

Epoch 29/40

1563/1563 [=====] - 150s 96ms/step - loss: 0.3265
- accuracy: 0.8841 - val_loss: 0.6167 - val_accuracy: 0.7971

Epoch 30/40

1563/1563 [=====] - 149s 96ms/step - loss: 0.3309
- accuracy: 0.8823 - val_loss: 0.6147 - val_accuracy: 0.7980

Epoch 31/40

1563/1563 [=====] - 150s 96ms/step - loss: 0.3216

In [537]:

```
model.fit(X_train, y_train, validation_data=(X_test, y_test), epochs=10, batch_size=32)
# Final evaluation of the model
scores = model.evaluate(X_test, y_test, verbose=0)
```

Epoch 1/10

```
1563/1563 [=====] - 154s 99ms/step - loss: 0.2813 - accuracy: 0.8991 - val_loss: 0.6202 - val_accuracy: 0.8006
```

Epoch 2/10

```
1563/1563 [=====] - 151s 97ms/step - loss: 0.2760 - accuracy: 0.9001 - val_loss: 0.6246 - val_accuracy: 0.8013
```

Epoch 3/10

```
1563/1563 [=====] - 151s 97ms/step - loss: 0.2733 - accuracy: 0.9017 - val_loss: 0.6243 - val_accuracy: 0.8025
```

Epoch 4/10

```
1563/1563 [=====] - 152s 97ms/step - loss: 0.2727 - accuracy: 0.9031 - val_loss: 0.6287 - val_accuracy: 0.8007
```

Epoch 5/10

```
1563/1563 [=====] - 149s 95ms/step - loss: 0.2675 - accuracy: 0.9034 - val_loss: 0.6294 - val_accuracy: 0.7993
```

Epoch 6/10

```
1563/1563 [=====] - 150s 96ms/step - loss: 0.2631 - accuracy: 0.9055 - val_loss: 0.6294 - val_accuracy: 0.7990
```

Epoch 7/10

```
1563/1563 [=====] - 154s 98ms/step - loss: 0.2630 - accuracy: 0.9050 - val_loss: 0.6224 - val_accuracy: 0.8017
```

Epoch 8/10

```
1563/1563 [=====] - 150s 96ms/step - loss: 0.2630 - accuracy: 0.9058 - val_loss: 0.6331 - val_accuracy: 0.8011
```

Epoch 9/10

```
1563/1563 [=====] - 152s 97ms/step - loss: 0.2595 - accuracy: 0.9068 - val_loss: 0.6266 - val_accuracy: 0.8006
```

Epoch 10/10

```
1563/1563 [=====] - 150s 96ms/step - loss: 0.2571 - accuracy: 0.9069 - val_loss: 0.6311 - val_accuracy: 0.8014
```

In [541]:

```
# second experiment
model_2 = Sequential()
model_2.add(Conv2D(16, (3, 3), activation='relu', strides=(1, 1), padding='same', input_shape=(28, 28, 1)))
model_2.add(Conv2D(32, (3, 3), activation='relu', strides=(1, 1), padding='same'))
model_2.add(Conv2D(64, (3, 3), activation='relu', strides=(1, 1), padding='same'))
model_2.add(MaxPooling2D((2, 2)))
model_2.add(Conv2D(16, (3, 3), activation='relu', strides=(1, 1), padding='same'))
model_2.add(Conv2D(32, (3, 3), activation='relu', strides=(1, 1), padding='same'))
model_2.add(Conv2D(64, (3, 3), activation='relu', strides=(1, 1), padding='same'))
model_2.add(MaxPooling2D((2, 2)))
model_2.add(Flatten())
model_2.add(Dense(256, activation='relu'))
model_2.add(Dropout(0.5))
model_2.add(Dense(128, activation='relu'))
model_2.add(Dense(64, activation='relu'))
model_2.add(Dense(64, activation='relu'))
model_2.add(Dense(10, activation='softmax'))
```


In [544]:

```
model_2.summary()
```

Model: "sequential_20"

Layer (type)	Output Shape	Param #
=====		
conv2d_46 (Conv2D)	(None, 32, 32, 16)	448
conv2d_47 (Conv2D)	(None, 32, 32, 32)	4640
conv2d_48 (Conv2D)	(None, 32, 32, 64)	18496
max_pooling2d_23 (MaxPooling)	(None, 16, 16, 64)	0
conv2d_49 (Conv2D)	(None, 16, 16, 16)	9232
conv2d_50 (Conv2D)	(None, 16, 16, 32)	4640
conv2d_51 (Conv2D)	(None, 16, 16, 64)	18496
max_pooling2d_24 (MaxPooling)	(None, 8, 8, 64)	0
flatten_9 (Flatten)	(None, 4096)	0
dense_19 (Dense)	(None, 256)	1048832
dropout_30 (Dropout)	(None, 256)	0
dense_20 (Dense)	(None, 128)	32896
dense_21 (Dense)	(None, 64)	8256
dense_22 (Dense)	(None, 64)	4160
dense_23 (Dense)	(None, 10)	650
=====		
Total params: 1,150,746		
Trainable params: 1,150,746		
Non-trainable params: 0		

In [550]:

```
# Compile model
epochs = 12
lr = 0.01
decay = lr/epochs
sgd = SGD(lr = lr, momentum= 1, decay=decay, nesterov=False,)
model_2.compile(loss='categorical_crossentropy', optimizer=sgd, metrics=['accuracy'])
model_2.fit(X_train, y_train, validation_data=(X_test, y_test), epochs=20, batch_size=32)
```

Epoch 1/20

1563/1563 [=====] - 106s 68ms/step - loss: nan - accuracy: 0.1021 - val_loss: nan - val_accuracy: 0.1000

Epoch 2/20

1563/1563 [=====] - 115s 73ms/step - loss: nan - accuracy: 0.1000 - val_loss: nan - val_accuracy: 0.1000

Epoch 3/20

1563/1563 [=====] - 114s 73ms/step - loss: nan - accuracy: 0.1000 - val_loss: nan - val_accuracy: 0.1000

Epoch 4/20

1563/1563 [=====] - 114s 73ms/step - loss: nan - accuracy: 0.1000 - val_loss: nan - val_accuracy: 0.1000

Epoch 5/20

1563/1563 [=====] - 112s 72ms/step - loss: nan - accuracy: 0.1000 - val_loss: nan - val_accuracy: 0.1000

Epoch 6/20

1563/1563 [=====] - 113s 72ms/step - loss: nan - accuracy: 0.1000 - val_loss: nan - val_accuracy: 0.1000

Epoch 7/20

1563/1563 [=====] - 115s 73ms/step - loss: nan - accuracy: 0.1000 - val_loss: nan - val_accuracy: 0.1000

In [553]:

```
model_2.fit(X_train, y_train, validation_data=(X_test, y_test), epochs=10, batch_size=32)
# Final evaluation of the model
scores = model_2.evaluate(X_test, y_test, verbose=0)
```

Epoch 1/10

1563/1563 [=====] - 117s 75ms/step - loss: nan - accuracy: 0.1000 - val_loss: nan - val_accuracy: 0.1000

Epoch 2/10

1563/1563 [=====] - 118s 75ms/step - loss: nan - accuracy: 0.1000 - val_loss: nan - val_accuracy: 0.1000

Epoch 3/10

1563/1563 [=====] - 117s 75ms/step - loss: nan - accuracy: 0.1000 - val_loss: nan - val_accuracy: 0.1000

Epoch 4/10

1563/1563 [=====] - 144s 92ms/step - loss: nan - accuracy: 0.1000 - val_loss: nan - val_accuracy: 0.1000

Epoch 5/10

1563/1563 [=====] - 117s 75ms/step - loss: nan - accuracy: 0.1000 - val_loss: nan - val_accuracy: 0.1000

Epoch 6/10

1563/1563 [=====] - 117s 75ms/step - loss: nan - accuracy: 0.1000 - val_loss: nan - val_accuracy: 0.1000

Epoch 7/10

1563/1563 [=====] - 119s 76ms/step - loss: nan - accuracy: 0.1000 - val_loss: nan - val_accuracy: 0.1000

Epoch 8/10

1563/1563 [=====] - 118s 75ms/step - loss: nan - accuracy: 0.1000 - val_loss: nan - val_accuracy: 0.1000

Epoch 9/10

1563/1563 [=====] - 117s 75ms/step - loss: nan - accuracy: 0.1000 - val_loss: nan - val_accuracy: 0.1000

Epoch 10/10

1563/1563 [=====] - 118s 75ms/step - loss: nan - accuracy: 0.1000 - val_loss: nan - val_accuracy: 0.1000